

A close-up photograph of a board game. The board is green with a pattern of dark grey and light grey circles. Various game pieces are scattered across the board: a green pawn, a yellow pawn, a black pawn, a white pawn, a red die, and a yellow die. The text "Heuristic Alpha-Beta Tree Search" is overlaid in the center in a white serif font.

Heuristic Alpha-Beta Tree Search

Heuristic Alpha–Beta Tree Search

- H-MINIMAX(s, d) for the heuristic minimax value of state s at search depth d :

$$\text{H-MINIMAX}(s, d) = \begin{cases} \text{EVAL}(s, \text{MAX}) & \text{if IS-CUTOFF}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if TO-MOVE}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if TO-MOVE}(s) = \text{MIN}. \end{cases}$$

Evaluation function

$$\text{UTILITY}(\text{loss}, p) \leq \text{EVAL}(s, p) \leq \text{UTILITY}(\text{win}, p)$$

- Good evaluation function?
 1. Computation must not take too long!
 2. The evaluation function should be strongly correlated with the actual chances of winning.

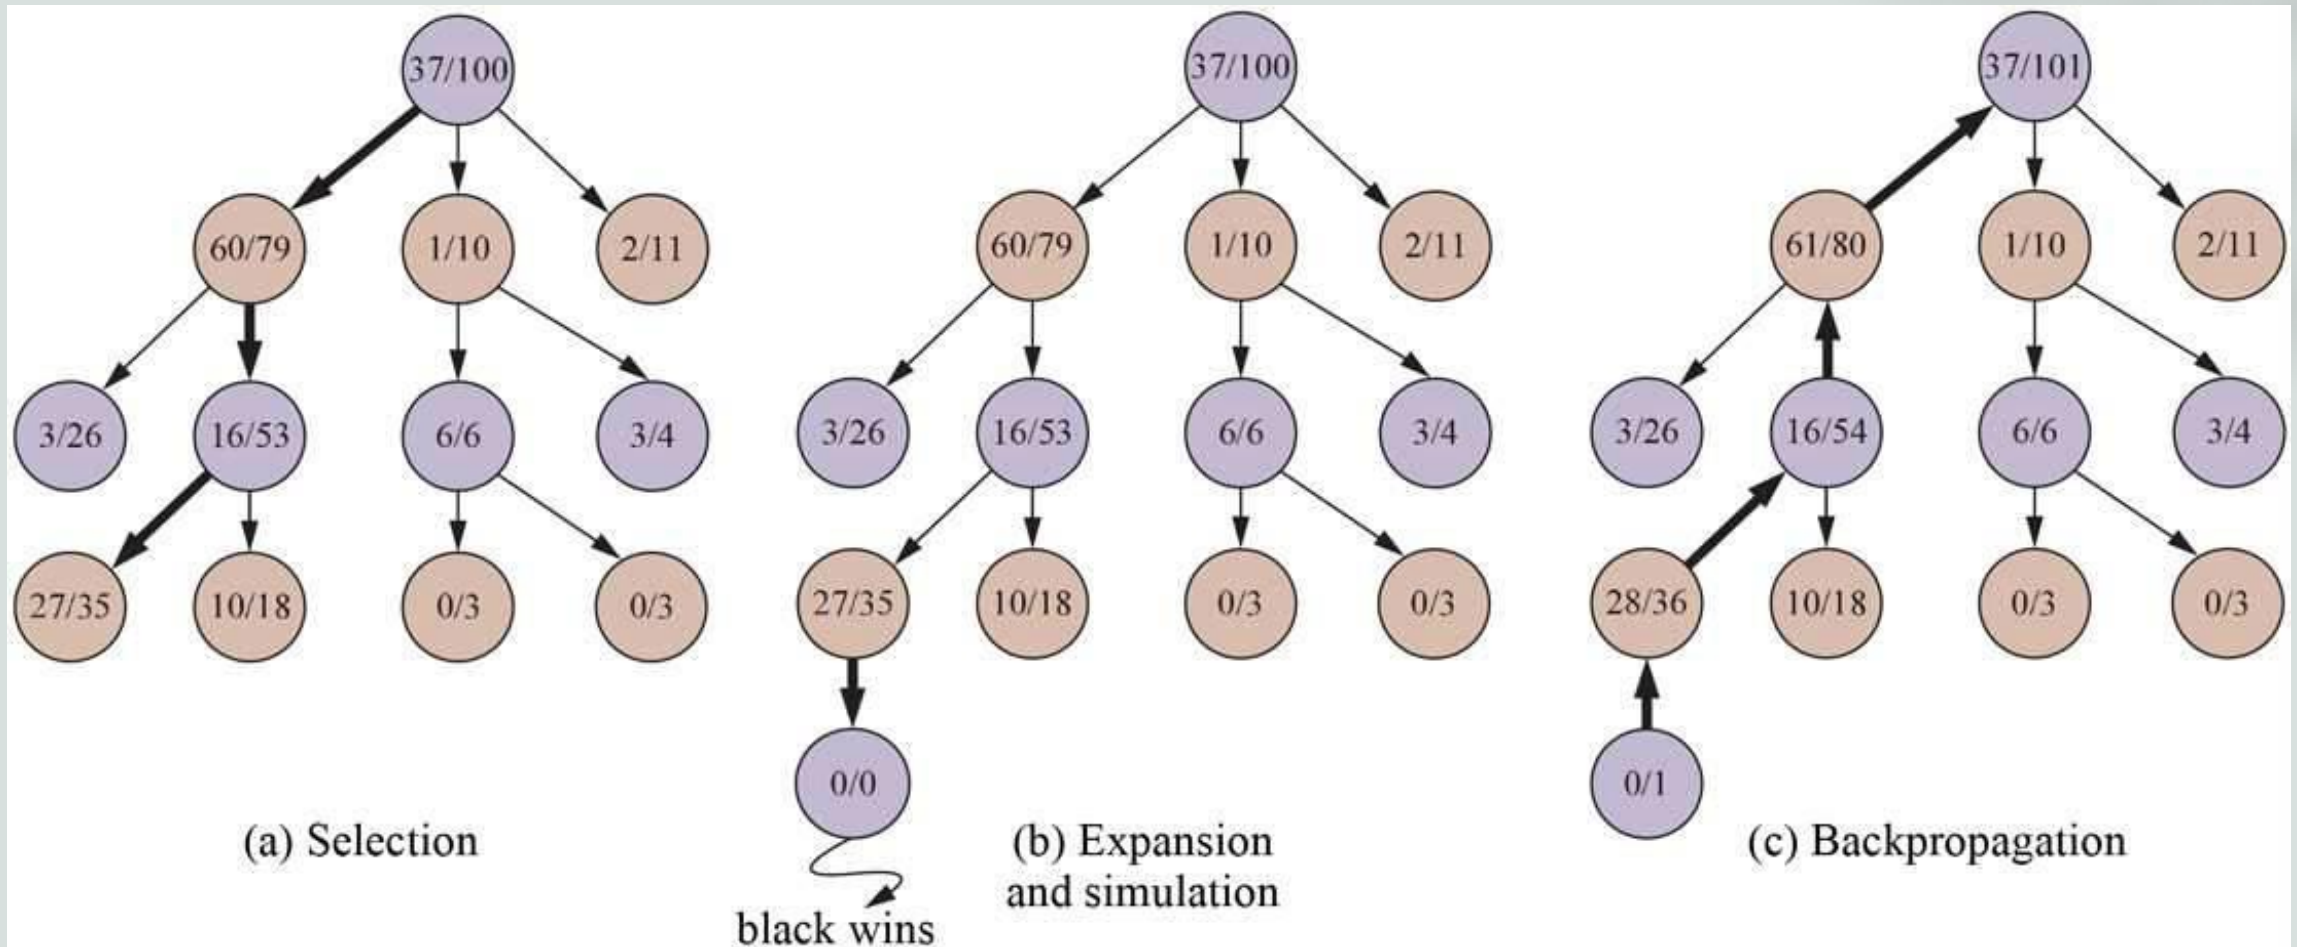
Monte Carlo tree search (MCTS)

- The basic MCTS strategy does not use a heuristic evaluation function.
- The value of a state is estimated as the average utility over a number of simulations of complete games starting from the state.
- A simulation (also called a playout or rollout) chooses moves first for one player, then for the other, repeating until a terminal position is reached.
- Playout policy that biases the moves towards good ones.
- From what positions do we start the playouts, and how many playouts do we allocate to each position?
- Pure Monte Carlo search, is to do N simulations starting from the current state of the game, and track which of the possible moves from the current position has the highest win percentage.

Monte Carlo tree search (MCTS)

- Selection policy
 - Exploration
 - Exploitation
- Selection
- Expansion
- Simulation
- Back-Propagatoin

Monte Carlo tree search (MCTS)



Selection Policy

- Upper confidence bounds applied to trees

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(\text{PARENT}(n))}{N(n)}}$$

- $U(n)$ is the total utility of all playouts that went through node n .
- $N(n)$ is the number of playouts through node n .
- $\text{PARENT}(n)$ is the parent node of n in the tree.
- C is a constant that balances exploitation and exploration.

Monte Carlo tree search (MCTS)

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action  
  tree  $\leftarrow$  NODE(state)  
  while IS-TIME-REMAINING() do  
    leaf  $\leftarrow$  SELECT(tree)  
    child  $\leftarrow$  EXPAND(leaf)  
    result  $\leftarrow$  SIMULATE(child)  
    BACK-PROPAGATE(result, child)  
  return the move in ACTIONS(state) whose node has highest number of playouts
```