

# Statistical Parsing

# Statistical parsing

- Over the last 12 years statistical parsing has succeeded wonderfully!
- NLP researchers have produced a range of (often free, open source) statistical parsers, which can parse *any sentence* and *often get most of it correct*
- These parsers are now a commodity component
- The parsers are still improving year-on-year.

# Statistical Parsing

- Basic idea
  - Start with a **treebank**
    - a collection of sentences with syntactic annotation, i.e., already-parsed sentences
  - Examine which parse trees occur frequently
  - Extract grammar rules corresponding to those parse trees, estimating the probability of the grammar rule based on its frequency
- That is, we'll have a CFG augmented with probabilities

# Treebanks

- Treebanks are corpora in which each sentence has been paired with a parse tree (presumably the right one).
- These are generally created
  - By first parsing the collection with an automatic parser
  - And then having human annotators correct each parse as necessary.
- This generally requires detailed annotation guidelines that provide a POS tagset, a grammar and instructions for how to deal with particular grammatical constructions.

# Penn Treebank

- Penn TreeBank is a widely used treebank.

- Most well known is the Wall Street Journal section of the Penn TreeBank.
  - 1 M words from the 1987-1989 Wall Street Journal.

```
( (S ( ' ' ' ' )
  (S-TPC-2
    (NP-SBJ-1 (PRP We) )
    (VP (MD would)
      (VP (VB have)
        (S
          (NP-SBJ (-NONE- *-1) )
          (VP (TO to)
            (VP (VB wait)
              (SBAR-TMP (IN until)
                (S
                  (NP-SBJ (PRP we) )
                  (VP (VBP have)
                    (VP (VBN collected)
                      (PP-CLR (IN on)
                        (NP (DT those)(NNS assets))))))))))
                )
              )
            )
          )
        )
      )
    )
  )
  ( , , ) ( ' ' ' ' )
  (NP-SBJ (PRP he) )
  (VP (VBD said)
    (S (-NONE- *T*-2) ))
  ( . . ) ) )
```

# Treebank Grammars



1

Treebanks implicitly define a grammar for the language covered in the treebank.



2

Simply take the local rules that make up the sub-trees in all the trees in the collection and you have a grammar.



3

Not complete, but if you have decent size corpus, you'll have a grammar with decent coverage.

# Treebank Grammars

- Such grammars tend to be very flat due to the fact that they tend to avoid recursion.
  - To ease the annotators burden
- For example, the Penn Treebank has 4500 different rules for VPs. Among them...


VP → VBD PP  
VP → VBD PP PP  
VP → VBD PP PP PP  
VP → VBD PP PP PP PP

# Heads in Trees



Finding heads in treebank trees is a task that arises frequently in many applications.

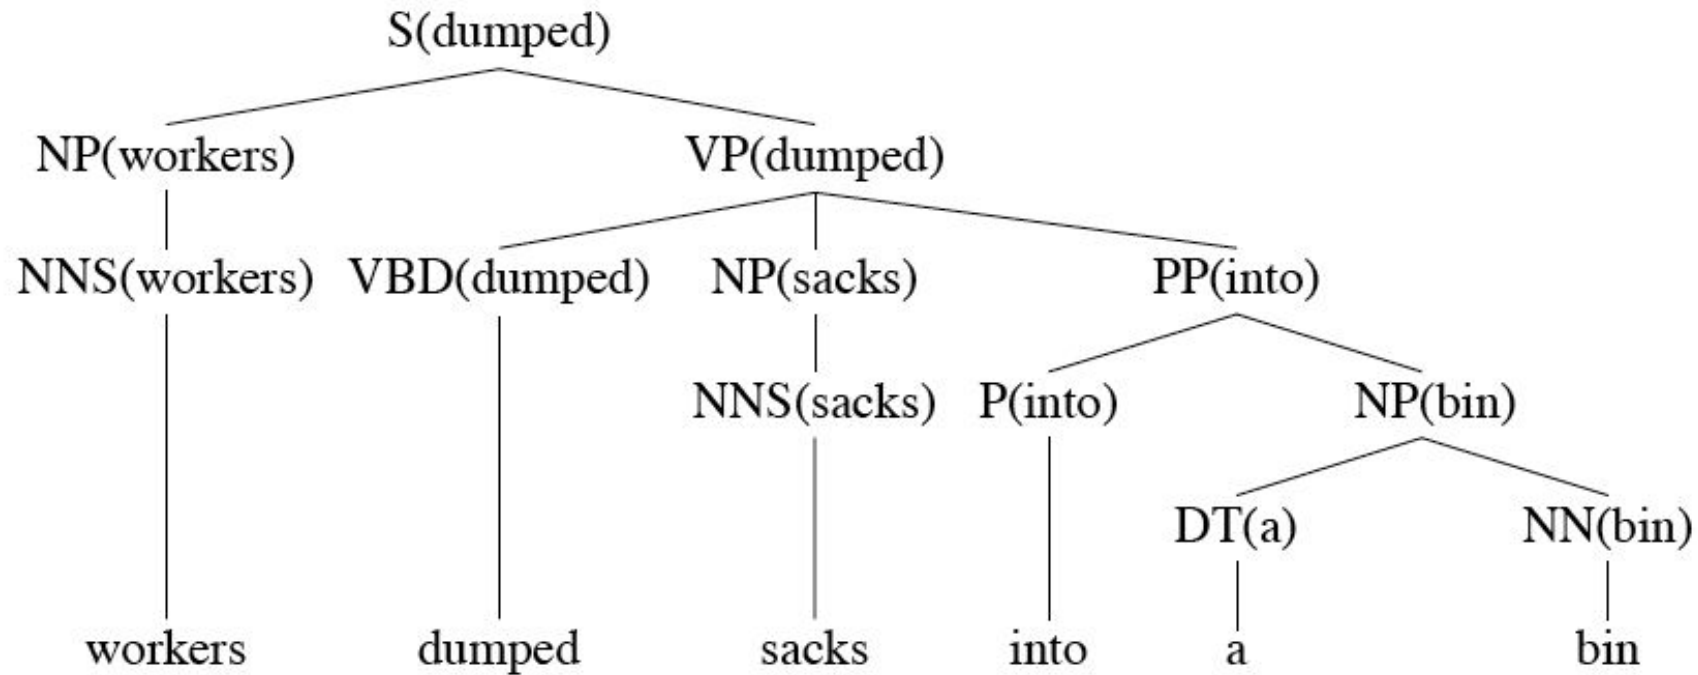
Particularly important in statistical parsing



We can visualize this task by annotating the nodes of a parse tree with the heads of each corresponding node.



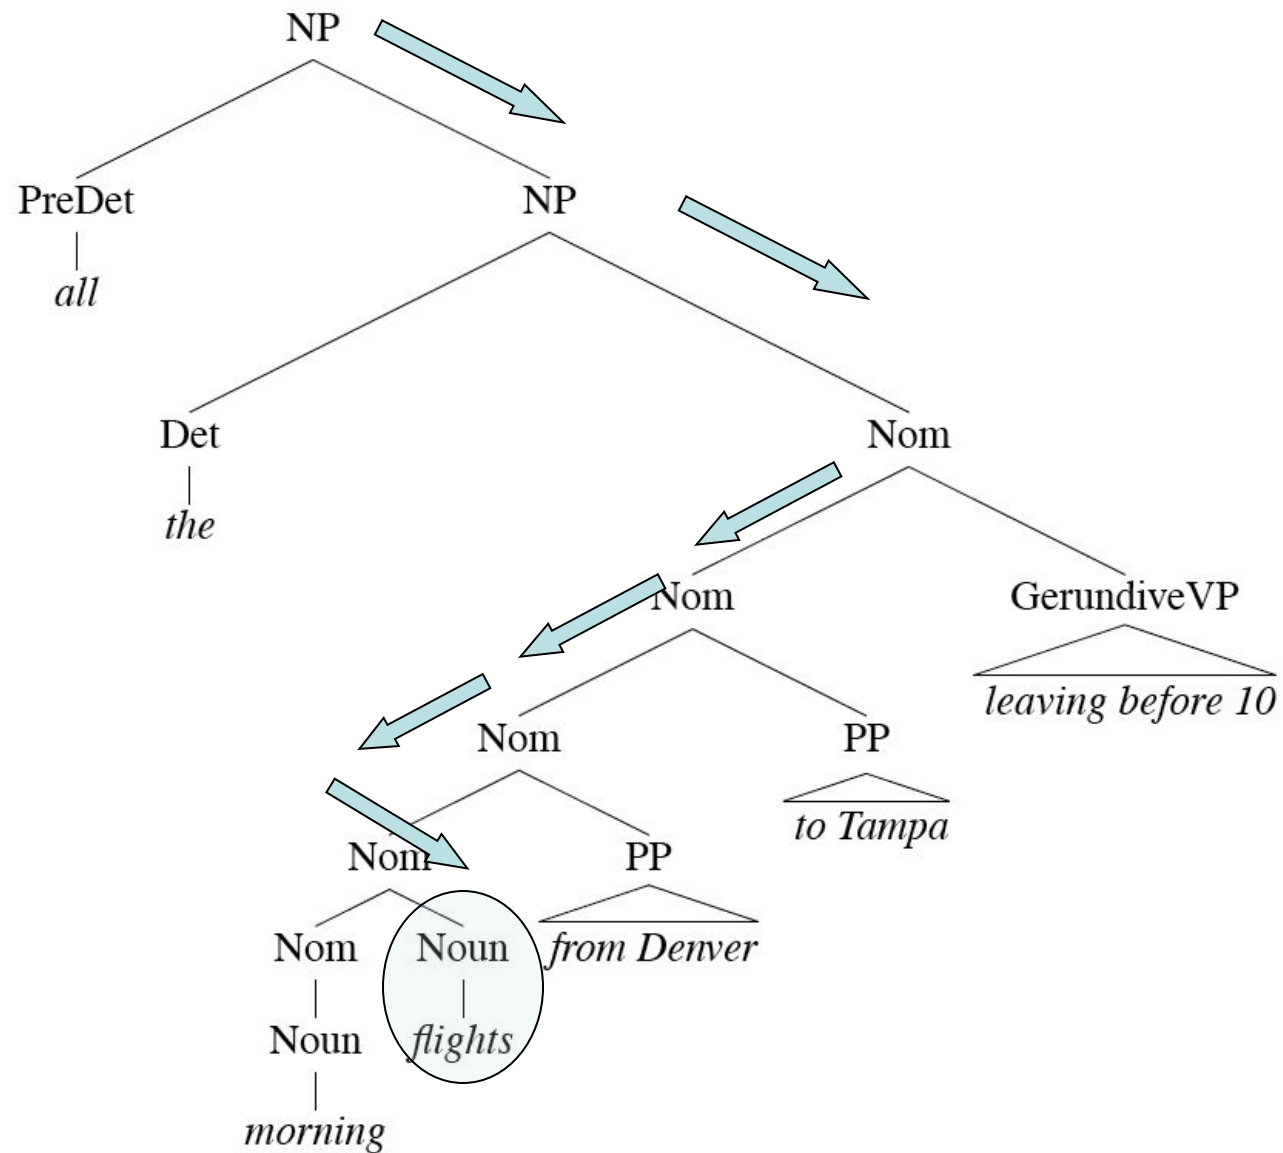
# Lexically Decorated Tree



# Head Finding

- The standard way to do head finding is to use a simple set of tree traversal rules specific to each non-terminal in the grammar.

# Noun Phrases



# Treebank Uses

- Treebanks (and headfinding) are particularly critical to the development of statistical parsers
- Also valuable to *Corpus Linguistics*
  - Investigating the empirical details of various constructions in a given language
    - How often do people use various constructions and in what contexts...
    - Do people ever say ...

# Probabilistic CFGs

## 1 The probabilistic model

- Assigning probabilities to parse trees

## 2 Training the model (Learning)

- Acquiring estimates for the probabilities specified by the model

## 3 Parsing with probabilities (Decoding)

- Given an input sentence, using the model to find the best (or n-best) tree for the input

# Rule Probabilities

- So... What's the probability of a rule?
- Start at the top...
  - A tree should have an **S** at the top. So given that we know we need an **S**, we can ask about the probability of each particular **S** rule in the grammar.
    - That is  $P(\text{particular rule} \mid S)$
- So in general we need

$$P(\alpha \rightarrow \beta \mid \alpha)$$

For each rule in the grammar

# Probability Model (1.1)

- The probability of a word sequence (sentence) is the probability of its tree in the unambiguous case.
- In the ambiguous case, it's the sum of the probabilities of the trees.
- Since we can use the probability of the tree(s) as a proxy for the probability of a sentence...
  - PCFGs give us an alternative to  $N$ -gram models as a kind of language model.

# Learning

- What if you have a corpus but don't have a treebank to get the counts from?
- Parse the corpus with a non-probabilistic grammar and collect counts for the rules that get used.
- Normalize
- Parse the corpus with a non-probabilistic grammar and collect counts for the rules that get used.
- Prorate the counts by the probability of the parse that it comes from.
- Normalize; update; iterate.



# Parsing (Decoding)

- So to get the best (most probable) parse for a given input
  1. Enumerate all the trees for a sentence
  2. Assign a probability to each using the model
  3. Return the argmax

# Formal Definition of PCFG

- A PCFG consists of
  - A set of terminals  $\{w_k\}$ ,  $k = 1, \dots, V$   
 $\{w_k\} = \{ \text{child, teddy, bear, played...} \}$
  - A set of non-terminals  $\{N_i\}$ ,  $i = 1, \dots, n$   
 $\{N_i\} = \{ \text{NP, VP, DT...} \}$
  - A designated start symbol  $N^1$
  - A set of rules  $\{N^i \rightarrow \zeta^j\}$ , where  $\zeta^j$  is a sequence of terminals & non-terminals  
 $\text{NP} \rightarrow \text{DT NN}$
  - A corresponding set of rule probabilities

# Rule Probabilities

- Rule probabilities are such that

$$\forall i \sum_j P(N^i \rightarrow \zeta^j) = 1$$

*E.g.*,  $P(\text{NP} \rightarrow \text{DT NN}) = 0.2$

$P(\text{NP} \rightarrow \text{NN}) = 0.5$

$P(\text{NP} \rightarrow \text{NP PP}) = 0.3$

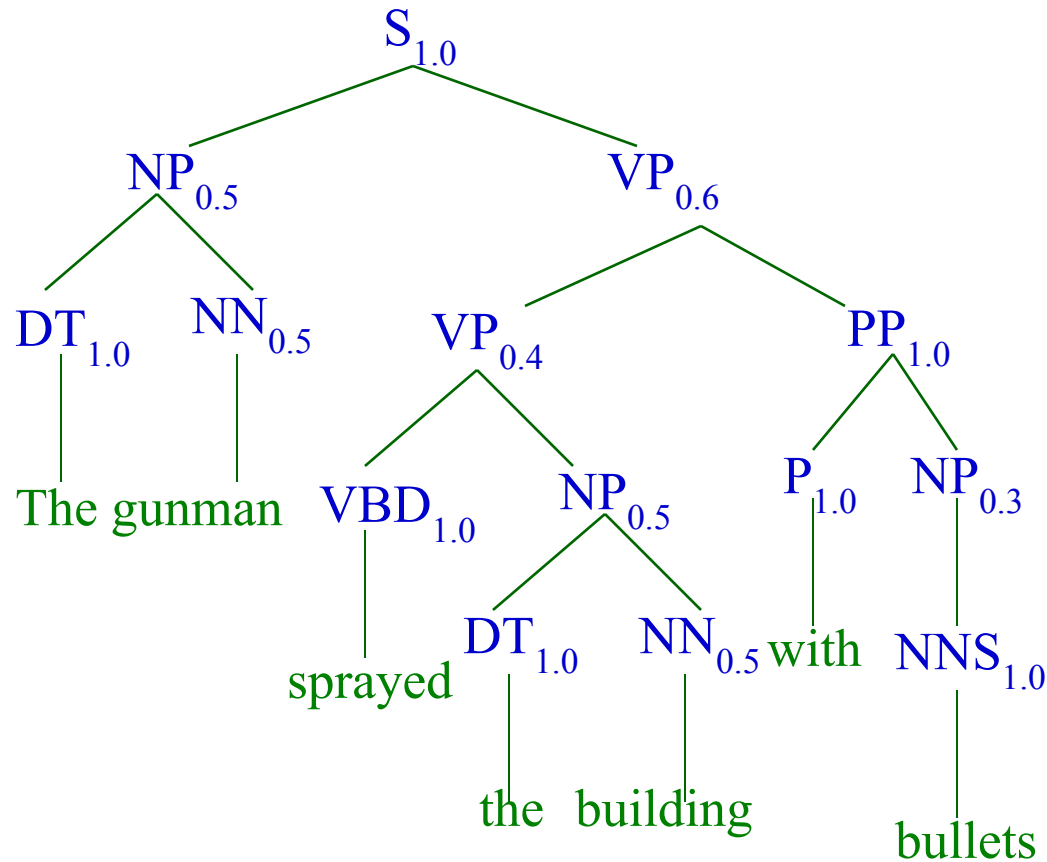
- $P(\text{NP} \rightarrow \text{DT NN}) = 0.2$ 
  - Means 20 % of the training data parses use the rule  $\text{NP} \rightarrow \text{DT NN}$

# Example PCFG Rules & Probabilities

- $S \rightarrow NP VP$  1.0
- $NP \rightarrow DT NN$  0.5
- $NP \rightarrow NNS$  0.3
- $NP \rightarrow NP PP$  0.2
- $PP \rightarrow P NP$  1.0
- $VP \rightarrow VP PP$  0.6
- $VP \rightarrow VBD NP$  0.4
- $DT \rightarrow the$  1.0
- $NN \rightarrow gunman$  0.5
- $NN \rightarrow building$  0.5
- $VBD \rightarrow sprayed$  1.0
- $NNS \rightarrow bullets$  1.0

# Example Parse $t_1$ .

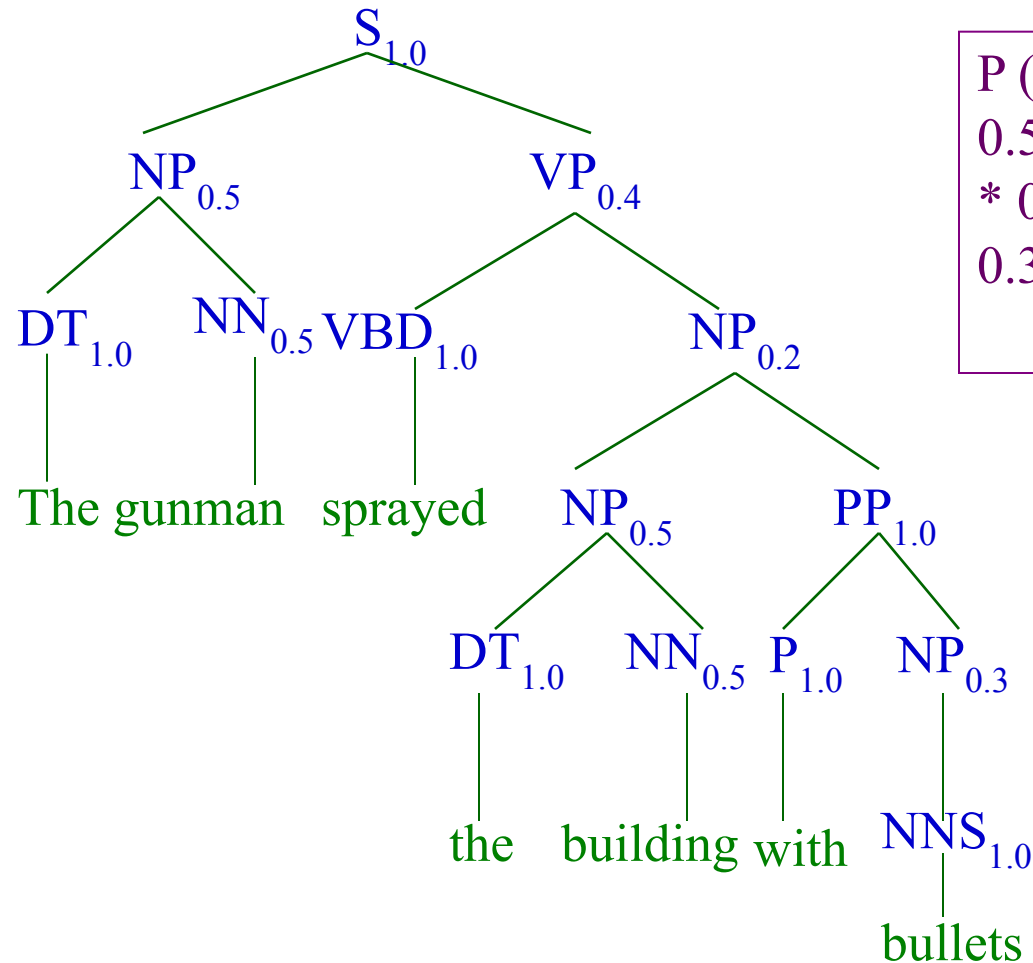
- The gunman sprayed the building with bullets.



$$\begin{aligned} P(t_1) &= 1.0 * \\ &0.5 * 1.0 * 0.5 * 0.6 * 0.4 * 1.0 \\ &* 0.5 * 1.0 * 0.5 * 1.0 * 1.0 * \\ &0.3 * 1.0 \\ &= \\ &0.00225 \end{aligned}$$

# Another Parse $t_2$

- The gunman sprayed the building with bullets.



$$\begin{aligned} P(t_2) &= 1.0 * \\ &0.5 * 1.0 * 0.5 * 0.4 * 1.0 * 0.2 \\ &* 0.5 * 1.0 * 0.5 * 1.0 * 1.0 * \\ &0.3 * 1.0 = 0.0015 \end{aligned}$$